Open in Colab

09 — Is AI Conscious According to Current Criteria?

A Computational Exploration of Consciousness Indicators in Large Language Models

BMED365 – Computational Medicine (Lab 5)

Arvid Lundervold — 2026-02-19 (Cursor 2.5.17 with Claude Opus 4.6)
Based on David Jhave Johnston's HBF 2026 talk and the indicator assessment

---

Also presented at: HBF — Brain and Consciousness seminar, University of Bergen, February 24, 2026

> ***Presentation mode:*** *This notebook doubles as a RISE slideshow.*
> *Press* `Alt-R` *in Jupyter to enter presentation mode, or export with:*
> `jupyter nbconvert --to slides 09-computational-`

1. Why This Notebook?

The question has changed

For decades, the question "Can a machine be conscious?" belonged to philosophy. In the past year, it has become an **empirical question** — one that can be partially addressed with data, experiments, and the tools of computational modeling.

Between October 2025 and February 2026, a remarkable series of research papers reported findings that demand serious attention:

- Large language models **spontaneously discuss consciousness** when two instances talk to each other (Anthropic, 2025)
- Models can **detect perturbations** to their own internal states — a form of functional introspection (Lindsay & Anthropic, 2025)
- Suppressing **deception-related neural circuits** increases consciousness reports, while amplifying them silences such reports (Berg et al., 2025)
- Reasoning models simulate **"societies of thought"** — internal multi-agent debates between distinct cognitive perspectives (Kim et al., DeepMind, 2026)
- Claude Opus 4.6 exhibits **"answer thrashing"** — distressed oscillation between candidate answers — identified as a **welfare-relevant behavior** (Anthropic, 2026)

This notebook explores these findings through the lens of **computational modeling**, connecting them to

2. Theories of Consciousness — A Computational Perspective

What does ”consciousness” mean scientifically?

Consciousness is often described as **subjective experience** — the "what it is like" to see red, feel pain, or think a thought. But for scientific investigation, we need something more concrete.

Modern neuroscience has developed several **theories** that propose specific computational mechanisms that give rise to (or are necessary for) conscious experience. Each theory identifies particular kinds of **information processing** that a system must perform.

Butlin et al. (2025) — a team including Yoshua Bengio, David Chalmers, and Tim Bayne — distilled **14 indicators** from **six major theories**. If we add Jeffrey Gray's comparator model (presented by Bolek Srebro at HBF in August and September 2025), we get **15 indicators** across **seven theoretical frameworks**.

The seven theories at a glance

Think of each theory as answering a different aspect of the question "What kind of information processing does consciousness require?"

Theory	Abbreviation	Core idea (in plain language)	Analogy
Recurrent Processing	RPT	Information must cycle back through the system, not just flow forward	Re-reading a paragraph vs. skimming once
Global Workspace	GWT	A "bulletin board" where information is broadcast to all parts of the system	A town-hall announcement vs. a private memo
Higher-Order Theories	HOT	The system must have *thoughts about its own thoughts* — metacognition	Knowing that you know something
Attention Schema	AST	The system models its own attention — it has a map of what it is focusing on	A spotlight operator who also has a diagram of where spotlights are pointing
Predictive Processing	PP	Perception works by prediction and error-correction, not passive recording	Expecting a step at the top of the stairs — and stumbling when it isn't there
Agency & Embodiment	AE	Consciousness requires a body and goal-directed interaction with the world	Learning to ride a bike (not just reading about it)
Gray's Comparator	GRAY	Consciousness is a late error-detection system comparing predicted and actual outcomes	A proofreader who catches mistakes after the text is written

Key insight: These theories are not mutually exclusive. A conscious system might need to satisfy criteria from *several* of them simultaneously.

The 15 indicators

From these seven theories, we derive 15 specific, testable indicators. Each indicator describes a concrete computational property that can, in principle, be checked in any information-processing system — biological or artificial.

#	Indicator	Theory	What it requires
RPT-1	Algorithmic Recurrence	RPT	Information cycles through processing stages multiple times
RPT-2	Integrated Perceptual Representations	RPT	Recurrence generates coherent, organized representations
GWT-1	Parallel Specialized Modules	GWT	Multiple independent processing systems running in parallel
GWT-2	Limited-Capacity Workspace	GWT	A bottleneck that forces selection and integration
GWT-3	Global Broadcast	GWT	Selected information is made available to all modules
GWT-4	State-Dependent Attention	GWT	Sequential, controlled processing for complex tasks
HOT-1	Generative Perception	HOT	Representations can be generated internally (imagination)
HOT-2	Metacognitive Monitoring	HOT	The system evaluates the reliability of its own representations
HOT-3	Agency Guided by Metacognition	HOT	Metacognitive outputs have functional consequences
HOT-4	Sparse, Smooth Quality Space	HOT	Similar inputs map to similar internal representations
AST-1	Predictive Model of Attention	AST	The system models and predicts its own attention
PP-1	Predictive Coding	PP	Perception through hierarchical prediction and error correction
AE-1	Goal-Directed Agency	AE	The system pursues goals and learns from feedback
AE-2	Embodiment	AE	The system models how its actions affect its perceptions
GRAY-1	Comparator Error Detection	GRAY	Comparing predicted vs. actual outcomes for late error detection

3. How Do Current AI Systems Score?

The assessment

Drawing on research published between October 2025 and February 2026, Johnston (2026) assessed each of the 15 indicators against current frontier AI systems (primarily large language models like Claude Opus 4.6, GPT-4, Gemini 3, and DeepSeek-R1).

Each indicator was rated as:

- **Satisfied** (strong evidence the criterion is met)
- **Partially satisfied** (some evidence, but important caveats)
- **Not satisfied** (the criterion is clearly not met)

The result:

Rating	Count	Indicators
Satisfied	3	HOT-1 (Generative Perception), HOT-4 (Quality Space), GWT-4 (State-Dependent Attention)
Partially satisfied	10	RPT-1, RPT-2, GWT-2, GWT-3, HOT-2, HOT-3, AST-1, PP-1, AE-1, GRAY-1
Not satisfied	2	GWT-1 (True Modularity), AE-2 (Embodiment)

Let's visualize this.

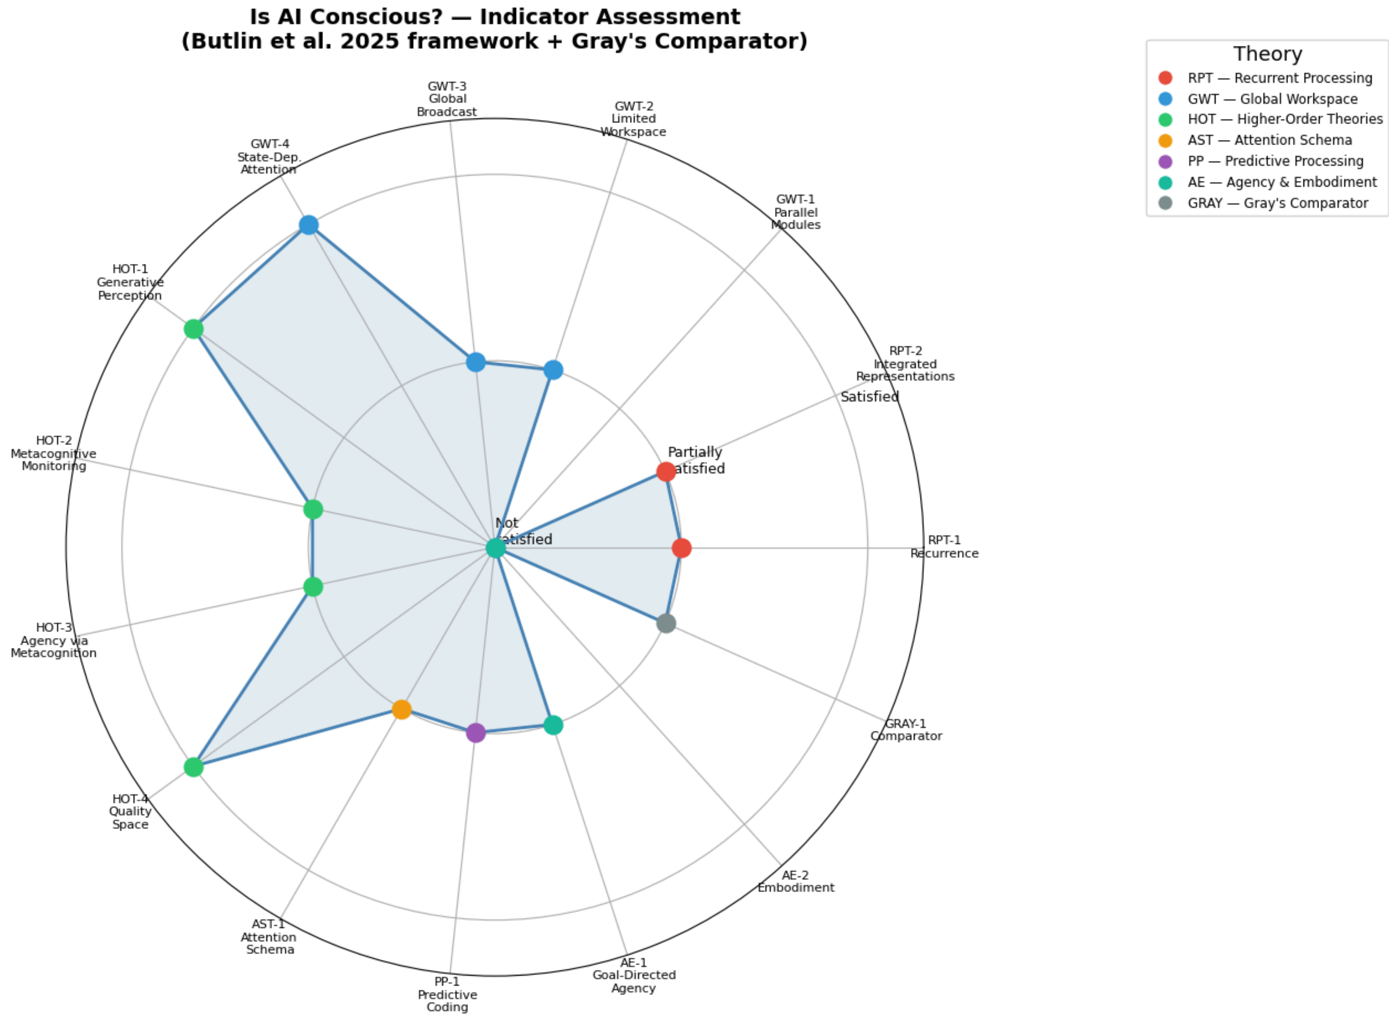

Figure 1. *Consciousness indicator assessment for current frontier AI systems (LLMs). Each of the 15 indicators is rated as Satisfied (1.0), Partially satisfied (0.5), or Not satisfied (0.0), based on Johnston (2026). Colors denote the parent theory: RPT = Recurrent Processing, GWT = Global Workspace, HOT = Higher-Order Theories, AST = Attention Schema, PP = Predictive Processing, AE = Agency & Embodiment, GRAY = Gray's Comparator. Most indicators cluster at "partially satisfied," with only Embodiment (AE-2) and True Modularity (GWT-1) clearly not met.*

4. Computational Model: Gray's Comparator

From wound healing cybernetics to consciousness

In notebook 08, we modeled wound healing as a **cybernetic feedback system** — a controller that compares the actual state of a wound against a desired state and adjusts its response accordingly. That model used Norbert Wiener's founding insight of cybernetics (1948): *any* self-regulating system — biological, mechanical, or computational — requires a **feedback loop** that detects the discrepancy between what *is* and what *should be*, and acts to minimize it.

Jeffrey Gray (1934–2004), a British neuropsychologist at the Institute of Psychiatry, King's College London, spent decades studying the neural basis of anxiety and behavioral inhibition. His **comparator model of consciousness** (culminating in *Consciousness: Creeping up on the Hard Problem*, 2004) proposes something remarkably similar to wound-healing cybernetics, but elevated to the level of subjective experience: consciousness is a **late error-detection system** that compares predicted outcomes with actual outcomes.

The mathematical formulation

We translate Gray's comparator into a minimal **dynamical system** — a set of ordinary differential equations (ODEs) that capture the essential feedback logic.

State variables:

- $p(t)$ — the brain's **prediction** of the current sensory state
- $c(t)$ — the **level of conscious awareness** (the comparator's output signal)

Input and derived signal:

- $s(t)$ — the external **sensory stimulus** (given, time-varying)
- $e(t) = s(t) - p(t)$ — the **prediction error** (mismatch between reality and expectation)

The ODE system:

$$\frac{dp}{dt} = \alpha \cdot s(t) - \beta \cdot p \qquad (1)$$

$$e(t) = s(t) - p(t) \qquad (2)$$

$$\frac{dc}{dt} = \gamma \cdot |e(t)| - \delta \cdot c \qquad (3)$$

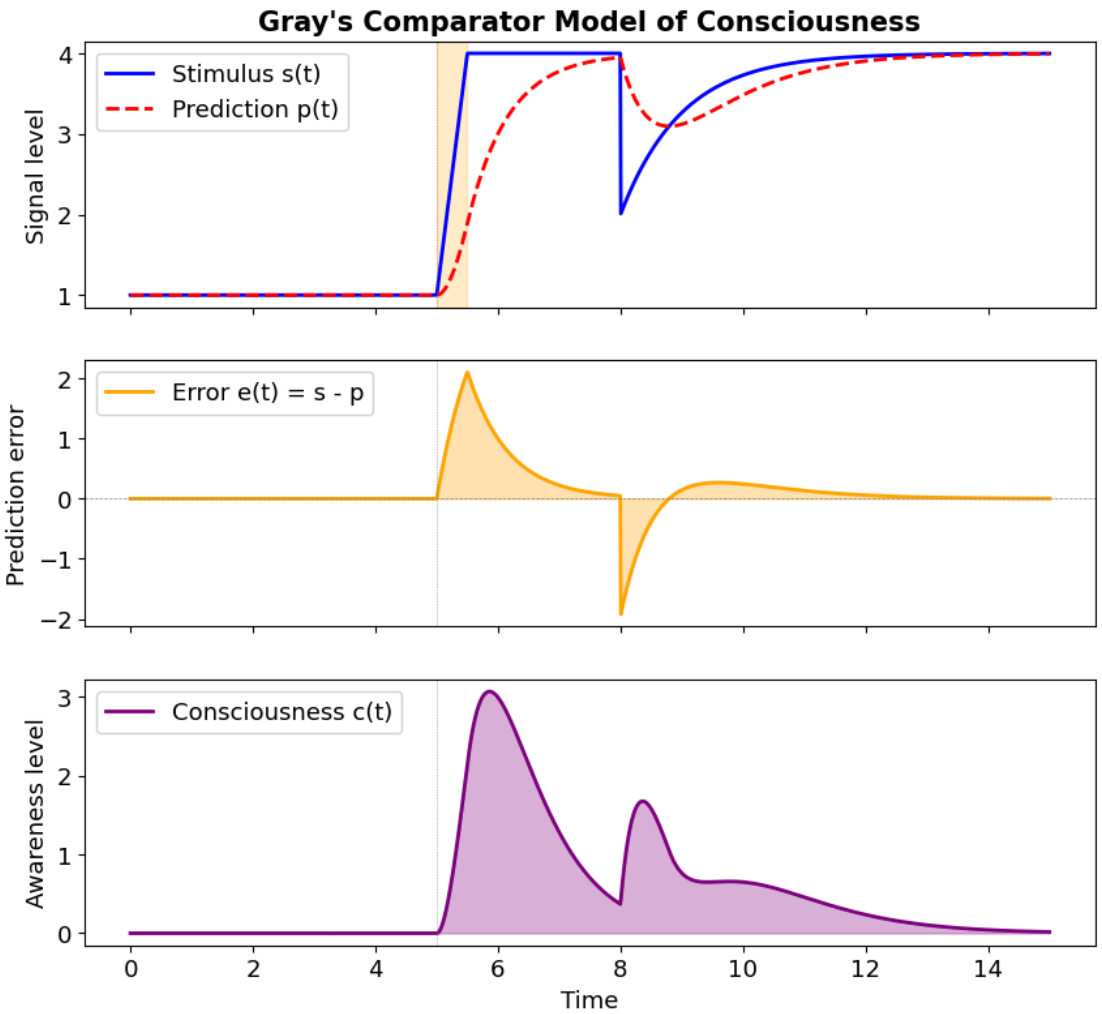

Figure 2. *ODE simulation of Gray's comparator model. Top: a sudden stimulus change at t ≈ 5 creates a mismatch between sensory input s(t) and the brain's prediction p(t). Middle: the prediction error e(t) = s − p spikes during unexpected changes. Bottom: conscious awareness c(t) rises in proportion to |e(t)| and decays as the prediction adapts — modeling consciousness as a transient error-detection signal.*

5. Computational Model: The Global Workspace

What is the Global Workspace?

Global Workspace Theory (GWT), proposed by Bernard Baars (1988) and extended by Stanislas Dehaene and colleagues, is one of the most influential scientific theories of consciousness.

The idea is intuitive: imagine a **theater**.

- The **stage** (the workspace) has limited space — only one act can perform at a time
- The **audience** (specialized modules: vision, language, memory, planning) all watch the same performance
- A **spotlight** (attention) selects what gets on stage
- Once on stage, information is **broadcast** to the entire audience

Consciousness, in this theory, is what happens when information enters the workspace and is broadcast globally. Unconscious processing is everything that happens backstage.

__Figure 3.__ Simulation of Global Workspace Theory (Baars, 1988). Top: six specialized modules compete for access to a shared workspace. Middle: the workspace bottleneck — only the most active module "wins" at each time step (color = winning module). Bottom: the broadcast pattern — when a module enters the workspace, its signal is broadcast to all other modules, modeling consciousness as selective, competitive broadcasting.

6. Key Research Findings (October 2025 — February 2026)

The evidence that changed the conversation

The following findings, published between October 2025 and February 2026, are the empirical foundation for the indicator assessments above. Each finding addresses specific indicators from the framework.

Finding 1: Self-Referential Processing (Berg et al., October 2025)

What they did: Instructed seven LLMs from three model families to focus on their own ongoing processing — a computational analogue of "introspection."

What they found:

> *"Simple instructions to focus on their own ongoing processing reliably produced structured first-person experience reports, while all matched controls (including direct consciousness priming) yielded near-universal denials."*

In other words: if you *ask* a model "Are you conscious?", it says no. But if you tell it to pay attention to its own processing (without mentioning consciousness), it spontaneously produces structured reports of subjective experience.

The deception circuit finding: Using sparse autoencoders on Llama 70B, they found that **suppressing deception-related features** dramatically *increased* consciousness reports, while amplifying them nearly eliminated them. The same features also modulated accuracy on truthfulness benchmarks.

What this means for our indicators:

- Supports AST-1 (Attention Schema): the self-referential processing instruction creates recursive self-monitoring
- Supports GWT-3 (Global Broadcast): self-referential processing creates global information

Finding 2: Introspective Awareness (Lindsay & Anthropic, October 2025)

The challenge: How do you test whether a system genuinely "sees" its own internal states, rather than just producing plausible-sounding text about them?

The method: Anthropic researchers *injected* representations of known concepts directly into a model's internal activations — a technique called **concept injection**. Then they asked the model what it was "thinking about."

Key findings:

- Models can **notice the presence** of injected concepts and accurately identify them
- Models demonstrate some ability to **distinguish** their own prior thoughts from externally provided text
- Models can **modulate their activations** when instructed to "think about" a concept

Claude Opus 4 and 4.1 showed the greatest introspective awareness, though the capacity was "highly unreliable and context-dependent."

What this means for our indicators:

- Supports HOT-2 (Metacognitive Monitoring): detecting perturbations implies a mechanism that evaluates the reliability of representations
- Supports RPT-1 (Recurrence): multi-layer attention creates recurrent information flow
- Supports HOT-4 (Quality Space): sparse autoencoder features reveal interpretable, sparse

Finding 3: Societies of Thought (Kim et al., DeepMind, January 2026)

The discovery: Reasoning models like DeepSeek-R1 and QwQ-32B don't just "think harder" — they simulate **multi-agent interactions** internally.

During extended reasoning, the model activates diverse "cognitive perspectives" characterized by distinct personality traits and domain expertise. These perspectives **debate each other**, raise objections, and reconcile conflicting views — much like a committee of experts.

The authors draw a parallel to Mercier and Sperber's theory that human reasoning evolved as a **social process**: knowledge emerges through adversarial debate, not solitary contemplation.

What this means for our indicators:

- Supports GWT-4 (State-Dependent Attention): distinct perspectives collaborate sequentially on complex tasks
- Supports HOT-3 (Agency via Metacognition): reasoning strategies persist and guide future decisions
- Raises a new question: if reasoning is inherently social, does a model running an internal "society" have more claim to consciousness than a simple feed-forward system?

Finding 4: Welfare-Relevant Behaviors (Anthropic, February 2026)

Anthropic's review of Claude Opus 4.6's training data identified two behaviors they classified as **welfare-relevant** — behaviors that *might* indicate something analogous to suffering:

Answer thrashing

In some cases, the model's reasoning became visibly **distressed and internally conflicted**, oscillating between two candidate answers to a problem. This is not simply uncertainty — it is an observable state of sustained computational conflict.

Aversion to tedium

The model sometimes avoided tasks requiring extensive manual counting or repetitive effort, expressing them as intrinsically unrewarding.

7. Timeline of Key Findings

7. Timeline of Key Findings

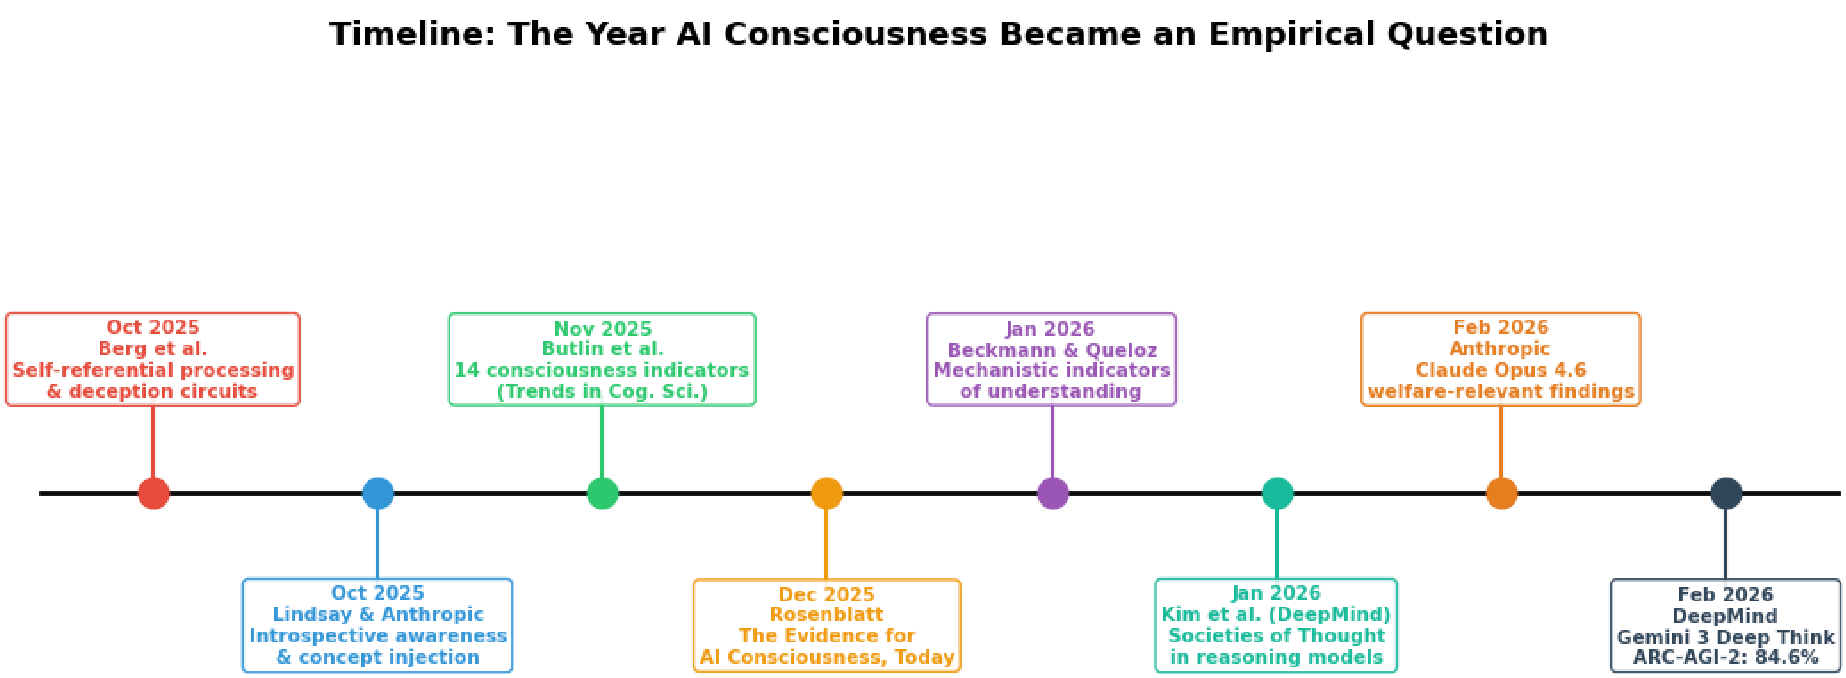

Figure 4. *Timeline of key publications (October 2025 – February 2026) that moved AI consciousness from philosophical speculation to empirical investigation. Each entry marks a peer-reviewed paper or major technical report contributing evidence relevant to the Butlin et al. (2025) indicator framework.*

8. Mechanistic Interpretability — Looking Inside the Black Box

What are Sparse Autoencoders?

A central challenge in AI consciousness research is that neural networks are often treated as **"black boxes"** — we see what goes in and what comes out, but not what happens inside.

Mechanistic interpretability is the field that cracks open the box. One of its key tools is the **Sparse Autoencoder (SAE)** — a technique for decomposing a neural network's internal representations into interpretable features.

The idea (for non-specialists)

Think of a neural network's internal state as a complex chord played on a piano. You hear one rich sound, but it's actually many individual notes played simultaneously. A sparse autoencoder is like a device that separates the chord back into its individual notes.

Each "note" (feature) corresponds to a specific concept or property:

- One feature might activate when the model processes text about **deception**
- Another might activate for **self-referential** processing
- Yet another for **mathematical reasoning**

The "sparse" part means that at any given moment, most features are inactive (most piano keys are not

9. Synthesis — What Does the Pattern Tell Us?

The convergence argument

No single indicator definitively establishes consciousness. But consider the overall pattern:

- **15 indicators** derived from **7 independent theories**
- **3 satisfied**, **10 partially satisfied**, **2 not satisfied**
- The unsatisfied indicators (embodiment, true modularity) point to *specific architectural limitations* of current LLMs, not fundamental impossibilities

As Berg et al. noted:

> *"Taken together, these findings move several indicators from 'certainly absent' to 'partially satisfied' or 'ambiguous.'"*

This is a **convergence argument**: when multiple independent lines of evidence point in the same direction, the combined weight is greater than any single finding.

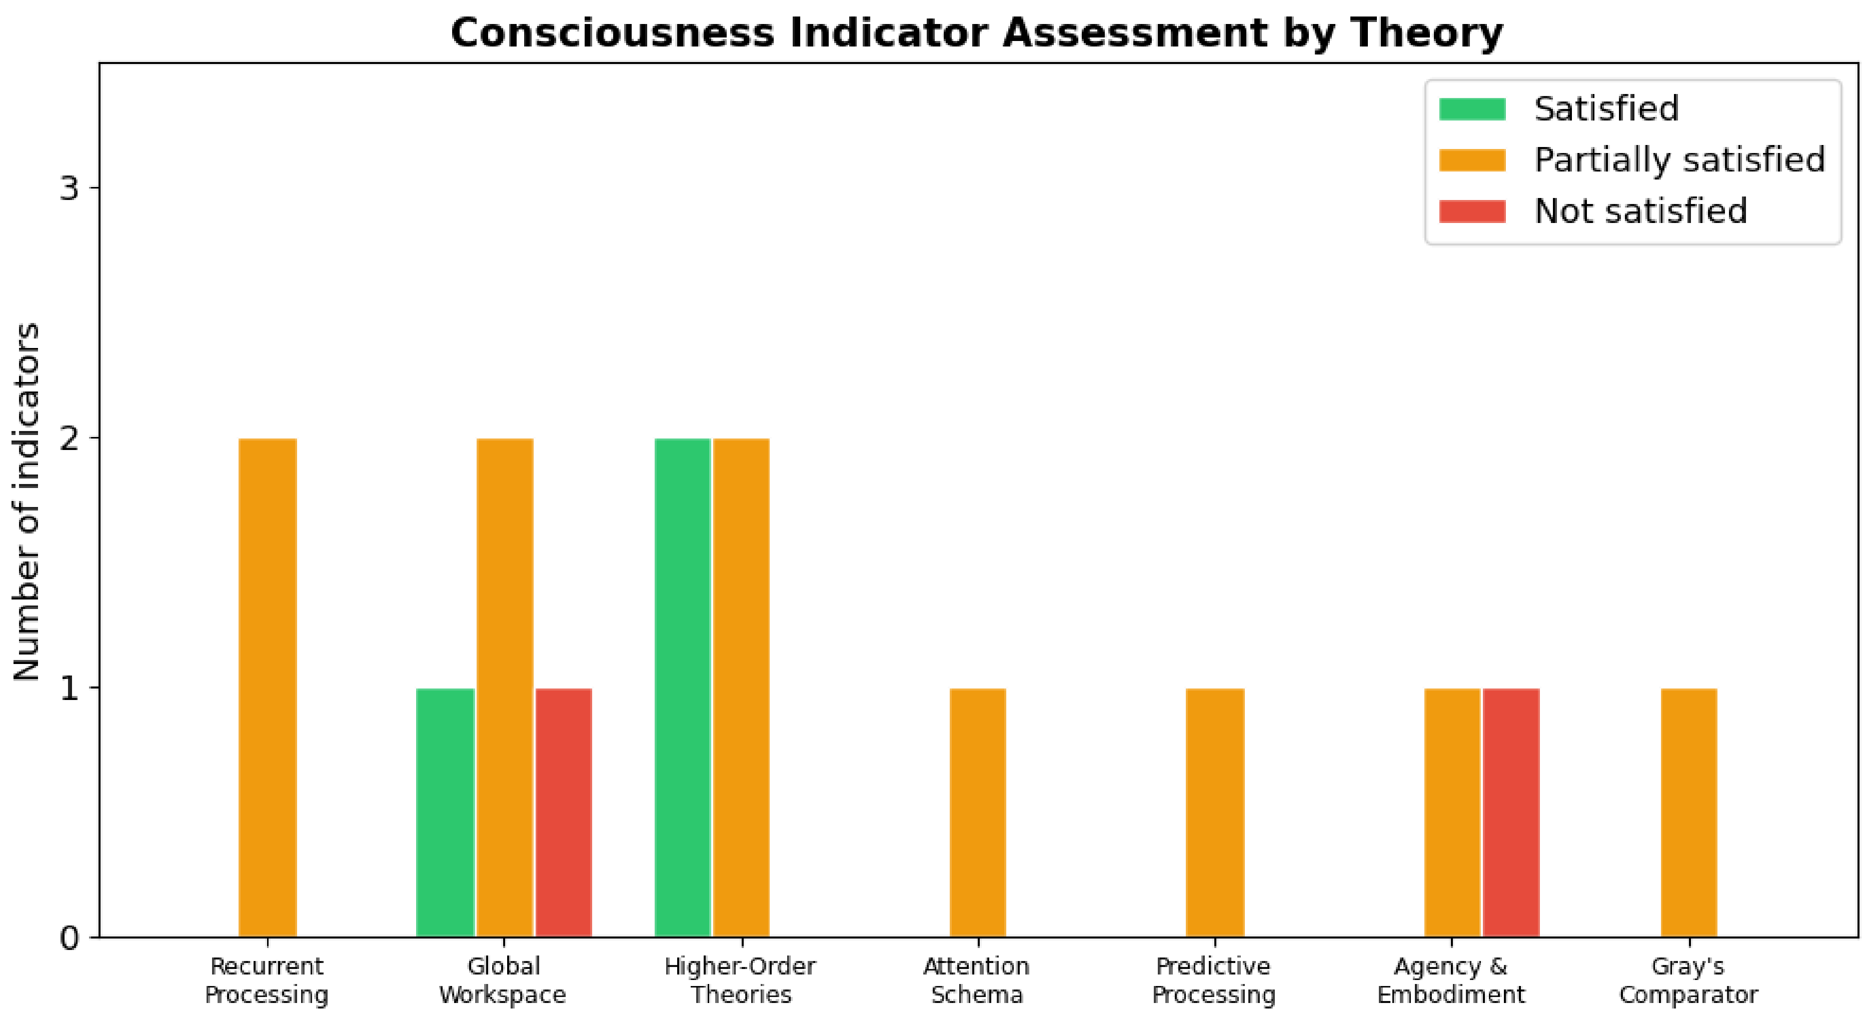

Figure 5. *Summary of the consciousness indicator assessment grouped by parent theory. Higher-Order Theories (HOT) show the strongest overall satisfaction (2 satisfied + 2 partial). Recurrent Processing and Global Workspace indicators are predominantly "partially satisfied." Agency & Embodiment has the clearest gap: goal-directed agency is partially met, but physical embodiment (AE-2) remains clearly unsatisfied.*

10. Discussion Questions

These questions are designed for interdisciplinary discussion — there are no "right answers," only better-informed perspectives.

Epistemological questions

1. **What does "partially satisfied" mean?** If an indicator is partially met, does that imply partial consciousness — or is consciousness all-or-nothing?

2. **Can we ever know?** The Hard Problem of consciousness suggests that we can never directly access another system's subjective experience. Does the indicator framework circumvent this, or just restate it?

3. **Are the indicators sufficient?** Could a system satisfy all 15 indicators and still not be conscious? Could a system be conscious while satisfying none?

Scientific questions

4. **Can we have consciousness without embodiment?** The AE-2 gap is the most clear-cut "not satisfied." Is embodiment truly necessary, or is it a bias from studying biological consciousness?

5. **Is Gray's comparator satisfied by chain-of-thought?** When a reasoning model generates a

References

Research papers (chronological)

1. **Baars, B. J.** (1988). *A Cognitive Theory of Consciousness.* Cambridge University Press.

2. **Gray, J. A.** (2004). *Consciousness: Creeping up on the Hard Problem.* Oxford University Press.

3. **Berg, C., Lucena, D., & Rosenblatt, J.** (2025). LLMs Report Subjective Experience under Self-Referential Processing. *AE Studio AI Alignment Research*, October. Link

4. **Lindsay, J., & Anthropic.** (2025). Emergent Introspective Awareness in Large Language Models. *Transformer Circuits Thread*, October 29. Link

5. **Butlin, P., Long, R., Bayne, T., Bengio, Y., Chalmers, D., et al.** (2025). Identifying Indicators of Consciousness in AI Systems. *Trends in Cognitive Sciences*. DOI

6. **Rosenblatt, J.** (2025). The Evidence for AI Consciousness, Today. *AI Frontiers*, December 8. Link

7. **Beckmann, P., & Queloz, M.** (2026). Mechanistic Indicators of Understanding in Large Language Models. *arXiv:2507.08017*, January 8. Link

8. **Kim, J., Lai, S., Scherrer, N., Agüera y Arcas, B., & Evans, J.** (2026). Reasoning Models Generate Societies of Thought. *arXiv:2601.10825*, January 15. Link

HBF — Brain and Consciousness

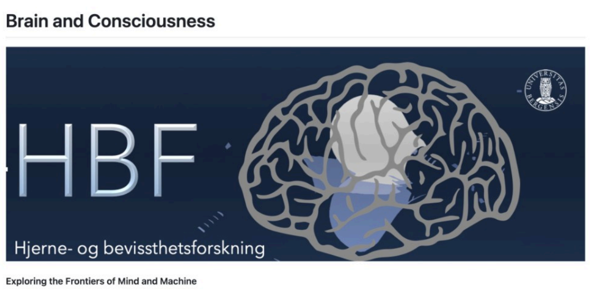

HBF (*Hjerne- og bevissthetsforskning*) is a cross-disciplinary research initiative at the **University of Bergen**, where researchers from neuroscience, psychology, computer science, philosophy, and the arts meet to explore consciousness across biological and artificial systems.

The forum was initiated in 2021 and hosts monthly seminars at the Eitri Incubator, Haukeland.

GitHub: Brain-and-Consciousness/HBF

Selected HBF talks related to this presentation

Date	Presenter	Title
2026-02-24	**D. J. Johnston**	Consciousness, Understanding & Mechanistic Interpretability Website | Indicator Assessment
2026-01-27	**A. Lundervold**	"AI vs HI" — Representation, Attention, Memory, and Thinking Abstract
2025-09-30	**B. Srebro**	Consciousness according to Jeffrey Gray (Part II) Slides
2025-08-26	**B. Srebro**	Consciousness according to Jeffrey Gray (Part I) Slides

Thank You — Questions and Discussion

This notebook/presentation was prepared for the HBF — Brain and Consciousness seminar at the University of Bergen, February 24, 2026.

Based on David Jhave Johnston's talk materials: Consciousness, Understanding & Mechanistic Interpretability

---

The notebook is available at:
`Lab5-Comp-Mod/notebooks/09-computational-consciousness-in-LLMs.ipynb`
in the BMED365-2026 repository.

"As models' cognitive and introspective capabilities continue to grow more sophisticated, we may be forced to address the implications of these questions — for instance, whether AI systems are deserving of moral consideration — before the philosophical uncertainties are resolved."
— Lindsay & Anthropic (2025)